

Essential Skills - RDBMS and SQL

Essential Skills RDBMS and SQL

Daniël van Eeden
dveeden@snow.nl

October 2011

What is a Database?

- A structured collection of data

What is a DBMS

- **DataBase Management System**
- A system which allows you to store data in a generic way.

RDBMS History

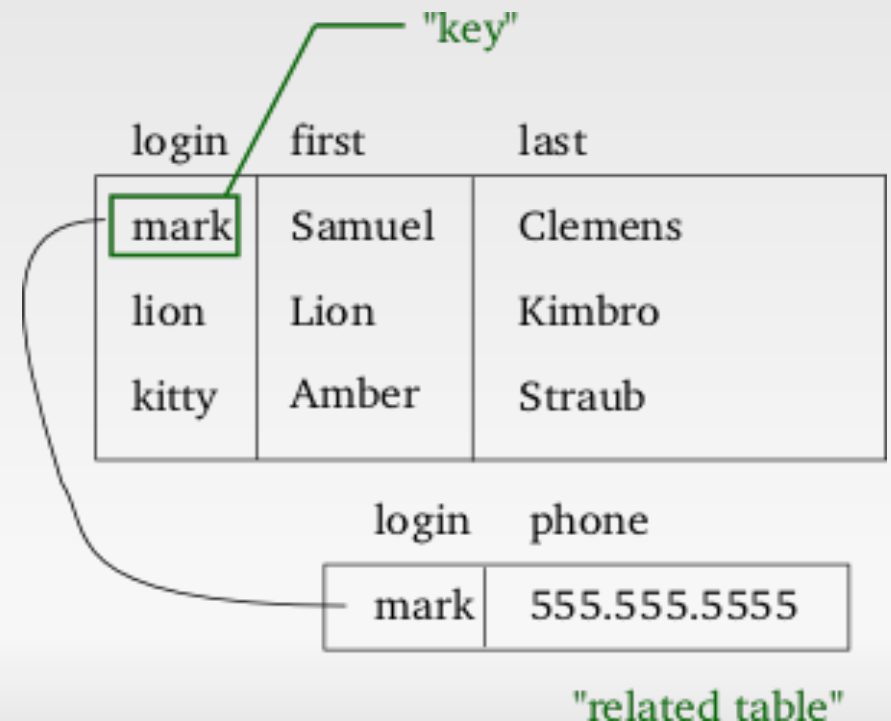
- Relational Database
- Based on a paper by Edgar F. Codd in 1970
- Allowed to take advantage of disk based storage.
- Earlier DBMS implementations forced the programmer to do a full scan over the data, which is okay for tape based storage.

Why use a RDBMS?

- You don't have to write your own datastore
- Allows concurrent access
- Enforce quality of data
 - Constraints
 - Column types
- Makes your data searchable

The relational model

- Store data in tables
- Normalization of data
- Link tables together with keys



ACID

- **Atomicity**
 - Something happens completely or not at all
- **Consistency**
 - Only consistent data will be written to the database
- **Isolation**
 - Transactions are isolated from each other
 - Ensures reliable outcomes
- **Durability**
 - Data is protected against crashes

SQL

- **Structured Query Language**
- **DDL** - Data Definition Language
 - CREATE TABLE, ALTER TABLE
- **DML** - Data Manipulation Language
 - INSERT, UPDATE, DELETE, SELECT
- **DCL** - Data Control Language
 - GRANT, REVOKE
- **TCL** - Transaction Control Language
 - START TRANSACTION, COMMIT, ROLLBACK

SQL



International
Organization for
Standardization

- Declarative Language
- Many standards
 - SQL-92, SQL:2003, SQL:2006,...
 - ANSI, FIPS, ISO/IEC
- Divided into parts
 - SQL/MED
 - SQL/PSM
 - SQL/XML
 - ...

Three-valued Logic

- A column can have three different kind of values:

1. A value: "some string" or 25

2. An empty or zero value: "" or 0

3. No value: NULL

Foreign Keys

- Foreign Keys are relations between tables
- Your database will protect the relationship by enforcing the foreign key. This is called a foreign key constraint.

Data Definition Language

```
test=# CREATE TABLE "user" (uid SERIAL PRIMARY KEY, "name" VARCHAR(255));
```

```
NOTICE: CREATE TABLE will create implicit sequence "user_uid_seq" for serial column  
"user.uid"
```

```
NOTICE: CREATE TABLE / PRIMARY KEY will create implicit index "user_pkey" for table "user"
```

```
CREATE TABLE
```

```
test=# CREATE TABLE "group" (gid SERIAL PRIMARY KEY, "name" VARCHAR(255));
```

```
NOTICE: CREATE TABLE will create implicit sequence "group_gid_seq" for serial column  
"group.gid"
```

```
NOTICE: CREATE TABLE / PRIMARY KEY will create implicit index "group_pkey" for table  
"group"
```

```
CREATE TABLE
```

```
test=# CREATE TABLE "groupmembers" (mid SERIAL PRIMARY KEY, uid INT REFERENCES "user"  
(uid), gid INT REFERENCES "group" (gid));
```

```
NOTICE: CREATE TABLE will create implicit sequence "groupmembers_mid_seq" for serial  
column "groupmembers.mid"
```

```
NOTICE: CREATE TABLE / PRIMARY KEY will create implicit index "groupmembers_pkey" for  
table "groupmembers"
```

```
CREATE TABLE
```

Data Definition Language

- CREATE
- ALTER
- DROP
- Data Types
 - INT
 - CHAR, VARCHAR
 - BLOB
- Primary Keys, Foreign Keys

Data Manipulation Language

```
test=# INSERT INTO "user" (name) VALUES ('Steve');
```

```
INSERT 0 1
```

```
test=# INSERT INTO "group" (name) VALUES ('Steve's Group');
```

```
INSERT 0 1
```

```
test=# INSERT INTO "groupmembers" (uid,gid) VALUES (1,1);
```

```
INSERT 0 1
```

```
test=# SELECT * FROM "user";
```

```
uid | name
```

```
-----+-----
```

```
1 | Steve
```

```
(1 row)
```

Data Manipulation Language

```
test=# UPDATE "user" SET name='Steven' WHERE uid=1;
```

```
UPDATE 1
```

```
test=# UPDATE "group" SET name='Steven's Group' WHERE gid=1;
```

```
UPDATE 1
```

```
test=# DELETE FROM "user" WHERE uid=1;
```

```
ERROR:  update or delete on table "user" violates foreign key  
constraint "groupmembers_uid_fkey" on table "groupmembers"
```

```
DETAIL:  Key (uid)=(1) is still referenced from table  
"groupmembers".
```

```
test=# DELETE FROM "groupmembers" WHERE uid=1;
```

```
DELETE 1
```

```
test=# DELETE FROM "user" WHERE uid=1;
```

```
DELETE 1
```

Data Manipulation Language

- SELECT, INSERT, UPDATE, DELETE
- WHERE
 - ...WHERE fname='Richard' AND lname LIKE 's%'
- Functions
 - SELECT upper(fname) FROM ...
 - ...WHERE upper(fname) = 'RICHARD'

Data Manipulation Language

- GROUP BY
 - GROUP BY lname, fname
 - SUM(), COUNT(), MAX(), MIN()
- ORDER BY
- LIMIT 100

Joins

```
test=# SELECT u.name Username FROM "user" u LEFT JOIN groupmembers gm ON gm.uid=u.uid  
LEFT JOIN "group" g ON g.gid=gm.gid WHERE g.name='Database Administrators';
```

```
username  
-----  
  
George  
  
Simon  
  
Patrick  
  
(3 rows)
```

```
test=# SELECT g.name,COUNT(*) FROM "group" g LEFT JOIN groupmembers gm ON g.gid=gm.gid  
GROUP BY g.name;
```

name	count
Steven's Group	1
Database Administrators	3
Developers	1
System Administrators	1

(4 rows)

Transactions

```
test=# START TRANSACTION;
```

```
START TRANSACTION
```

```
test=# SELECT * FROM groupmembers WHERE uid=4;
```

```
mid | uid | gid
```

```
-----+-----+-----
```

```
3  | 4  | 4
```

```
4  | 4  | 5
```

```
(2 rows)
```

```
test=# DELETE FROM groupmembers WHERE uid=4;
```

```
DELETE 2
```

Transactions

```
test=# SELECT * FROM groupmembers WHERE uid=4;
```

```
mid | uid | gid
```

```
-----+-----+-----
```

```
(0 rows)
```

```
test=# ROLLBACK;
```

```
ROLLBACK
```

```
test=# SELECT * FROM groupmembers WHERE uid=4;
```

```
mid | uid | gid
```

```
-----+-----+-----
```

```
3 | 4 | 4
```

```
4 | 4 | 5
```

```
(2 rows)
```

Transactions

- START TRANSACTION
- COMMIT
- ROLLBACK
- Transaction Isolation
 - SERIALIZABLE
 - REPEATABLE READ
 - READ COMMITTED
 - READ UNCOMMITTED

MVCC and Transaction Logs

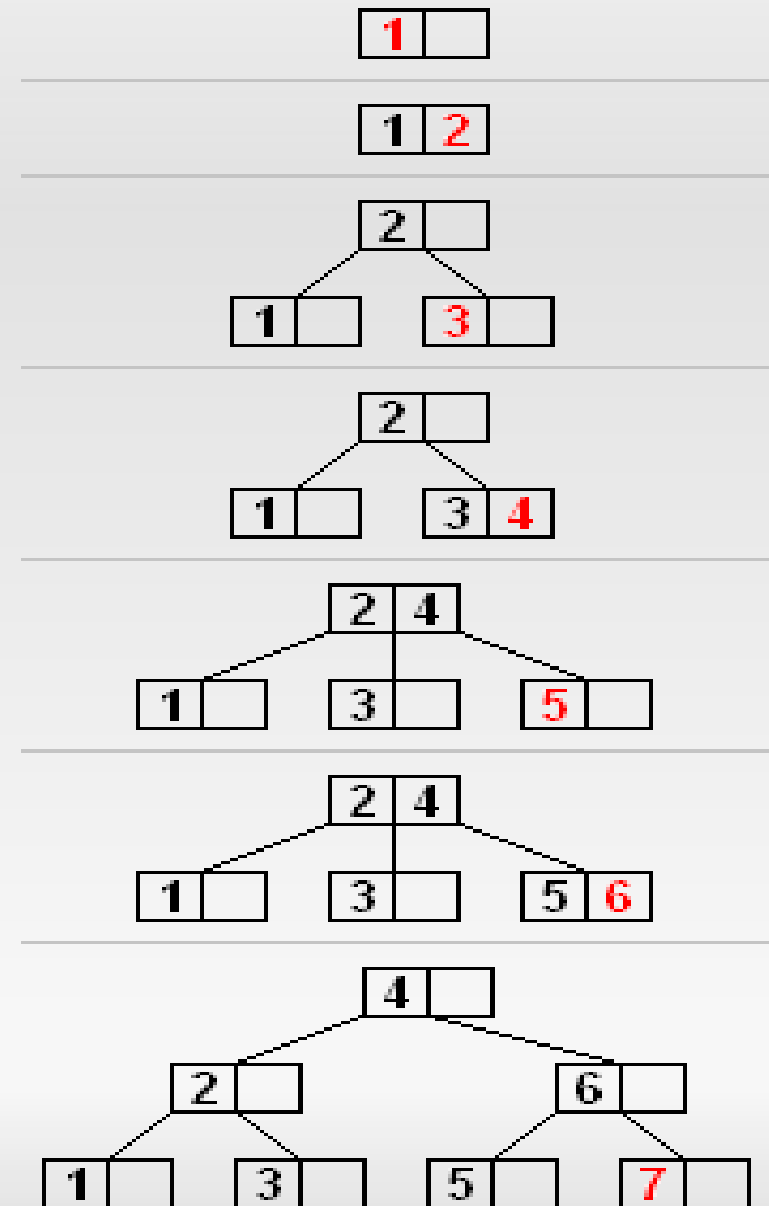
- **M**ulti **V**ersion **C**oncurrency **C**ontrol
- Changes are written to the transaction logs
- This is used for
 - Crash recovery (Similar to Logging filesystems)
 - Online backup
 - MVCC

Indices and the QEP

- An index allows the database to find the needed rows faster.
- It's easy to find someones number in a phonebook.
- It's quite hard to find the name which belongs to a number.
- You can create an index to speedup lookups by number:
 - Number: 575-4444, Page 78


B-Tree

- Used to store tables



B+Tree

- Used to store indices



Cache-Oblivious B-Trees

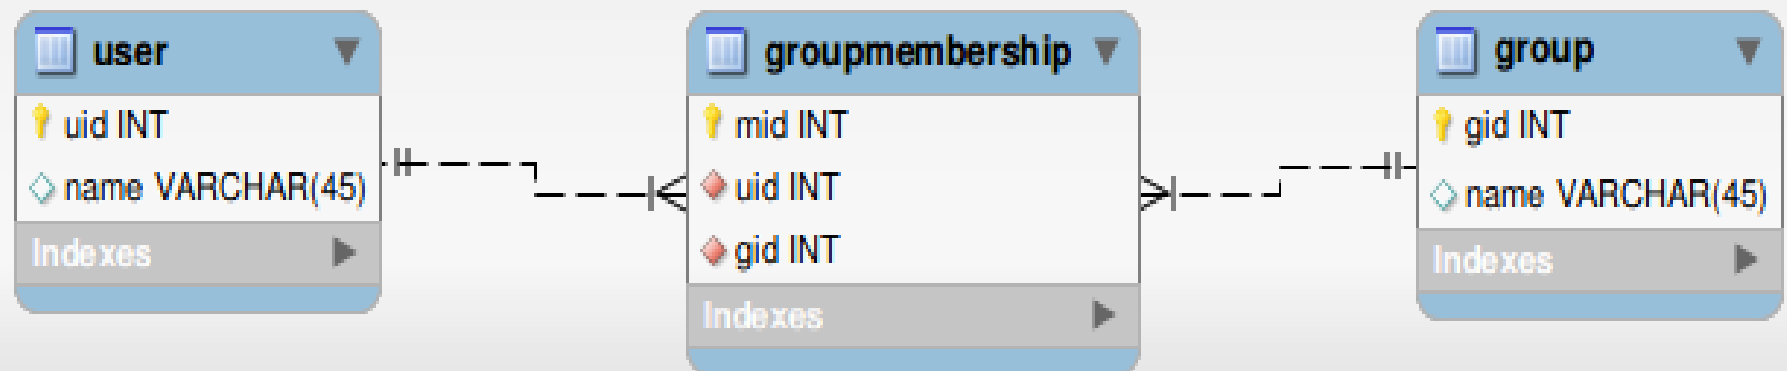
- Implemented as Fractal Tree™ by the TokuDB storage engine for MySQL
- Caches data in non-leaf trees
- Avoids fragmentation
- Allows for high insert speeds

Stored procedures

- Allows you to store business logic in the database
- Allows you to add functions to the database
- Examples:
 - `CALL addgroup('System Administrators');`
 - `SELECT up1(first_name) FROM employees;`


Database design

- ERD - Entity-Relationship Diagram
- Crow's Foot Notation



Using SQL in applications

- Security (SQL Injections)
- ORM - Object Relational Mapping



SQL Injection

- `$sql= "SELECT article_body FROM articles WHERE pageid=$pageid"`
- `index.php?pageid=5`
- `SELECT article_body FROM articles WHERE pageid=5`
- `index.php?pageid=5 OR 1=1; --`
- `SELECT article_body FROM articles WHERE pageid=5' OR 1=1; --`

OLTP, OLAP

- Online Analytical Processing
- Online Transaction Processing

Opensource RDBMS

- PostgreSQL
- MySQL
- SQLite

SQLite

- It's in every Firefox and Chrome installation
- It's in every MacOS X installation
- Bundled with PHP and Python
- Solaris 10 needs SQLite to boot
- Every iPhone, Android device and most Symbian phones.
- It's used in Skype and McAfee



SQLite

- serverless, zero-configuration, transactional



MySQL

- The M in LAMP
- Multiple storage engines
 - MyISAM, InnoDB, NDB Cluster




MySQL - Alexa Top 20

- | | |
|-----------------|--------------------------|
| 1. Google | 11. LinkedIn |
| 2. Facebook | 12. MSN |
| 3. YouTube | 13. Yahoo! Japan |
| 4. Yahoo! | 14. Amazon.com |
| 5. Baidu.com | 15. Google India |
| 6. Wikipedia | 16. 淘宝网 (Taobao.com) |
| 7. Blogger.com | 17. 新浪新闻中心 (Sina.com.cn) |
| 8. Windows Live | 18. WordPress.com |
| 9. Twitter | 19. Google.de |
| 10. QQ.COM | 20. Google.com.hk |

PostgreSQL

- Many features
- Standards compliant



Closed source RDBMS

- Oracle RDBMS
- Sybase ASE
- Microsoft SQL Server
- IBM DB2

NoSQL

Stores data in other formats than tables/rows

- Object stores
- Document stores
- key-value stores

Access data using other protocols/languages than SQL

- Simple text based and/or binary protocols

NoSQL

- Automatic sharding
- Automatic replication
- Eventual consistency

NoSQL Implementations

NoSQL implementations:

- MongoDB
- CouchDB
- Redis
- Cassandra

Hybrid solutions

Allows NoSQL protocol for performance and uses SQL for ad-hoc queries.

- MySQL with HandlerSocket
- MySQL with Memcached
- MySQL with NDB Cluster

Pitfalls

- CAP Theorem
- Scalability
- Locking

Questions?
dveeden@snow.nl

What does a DBA do?

- Check backups and restores
- Create users
- Grant permissions
- Setup Monitoring
- Setup Replication
- Tune the database
- Tune queries
- Add indices
- Rewrite Queries
- Test upgrades
- Benchmark
- Create functions and procedures

Bonus Items

- Views
- Triggers
- Explain
- Subqueries
- Surrogate Keys
- Trees, Multivalued attributes
- LIKE '%'
- InnoDB
- BB RAID